

File management and manipulation

Операционни системи, ФМИ, 2019/2020

Filesystem Hierarchy Standard

- Filesystem standard (FHS)
 - Guiding principles for each area of filesystem
 - Predictable location of files and directories
 - Provides uniformity across multiple Linux distributions
- The Linux Standards Base
 - Aims to allow Linux binaries to run unmodified on multiple Linux distributions
 - Specifies system and library interfaces and environment
 - Incorporates the FHS

Navigating the filesystem

- *absolute* vs. *relative* addressing
- changing and displaying directories (cd, pwd)
- cd (without parameters)
- cd ~george
- cd ~
- cd -
- . and ..

Displaying directory contents

- *human-readable*
- `ls`
- `ls -a` - show all files (including *.hidden* files)
- `ls -l` - long listings
 - `ls -lh`
- `ls -d` - show directories, not contents
- `touch foo`
- `mkdir bar`



File {group,}ownership

- each file is owned by a specific UID and GID
- `chown` - change the user (UID) ownership
 - only root can change ownership to another user
 - `chown foo:bar`
- `chgrp` - modify just the group (GID) ownership
 - `chown :bar`
- newly created files will usually be given GID ownership based on the current active group of the person who creates the file

File permissions

- type of file
 - - - regular file
 - b - block special file
 - c - character special file
 - d - directory
 - l - symbolic link
 - p - FIFO (named pipe)
 - s - socket
- permission sets
 - user (owner)
 - group (group owner)
 - everyone else (other)
 - symbolic representation `rwxr-xr-x`
 - numeric representation 0755

File permissions (cont.)



- r - 100b - 4 - Read
- w - 010b - 2 - Write
- x - 001b - 1 - Execute
- v

Special permissions

- Set UID upon execution (SUID)
- Set GID upon execution (SGID)
- sticky bit
- different behavior for files and directories

The diagram illustrates the conversion of octal permissions to symbolic notation. The top row shows octal permissions: 4 2 1 4 2 1 4 2 1. Below this is the symbolic notation: **rwxrwxrwx**. Three arrows point from the octal digits 4, 2, and 1 to the first three characters of the symbolic notation. The first arrow points from 4 to 'r', the second from 2 to 'w', and the third from 1 to 'x'. The word **SUID** is written in red above the first arrow. Below the symbolic notation is another set of permissions: **rwsrwxrwx**. The first three characters 'rws' are underlined, and the word **USER** is written in red below the underline.

4 2 1 4 2 1 4 2 1
rwxrwxrwx
SUID
rwsrwxrwx
USER

Special permissions (cont.)

- SUID and SGID on files
 - an executable with the SUID bit set runs with the security context of the user who owns it, regardless of the executing user
 - SGID
- SGID on directories
 - files or sub-directories created within that directory inherit the group ownership of the SGID directory
- Sticky Bit on directories
 - normally in a directory that is world writable, users can delete each other's files. Setting the sticky bit overrides this behavior

Changing file permissions

- `chmod`
 - numeric notation `chmod 0664 foo.txt`
 - symbolic notation `chmod u=rw,g=rw,o=r foo.txt`
 - `+`, `-`, `=`
- `chmod -R`

umask

- Default permissions for newly created filesystem objects
 - files 666
 - directories 777
- `umask`
 - defines what permissions to **withhold** from the default permissions
 - display or change umask
 - usually set in the user or system shell dot files


Directory and file manipulation

- `mkdir foo`
 - `mkdir -p foo/bar`
 - `mkdir -m`
- `rmdir`
- `cp`
- `mv`
- `rm`
- `touch - mtime/ctime`

UNIX filesystem structure


- blocks
- inodes
 - permissions
 - access time, modification time, inode change time
 - owner
 - group
 - size in bytes
 - occupied blocks
 - link count (names of the inode)
 - inode number
- directories (are files that) hold filenames and inodes
- superblock contains filesystem parameters (how many inodes, etc.)

inode pointer structure



Filesystem hard links


- a directory entry that references the same inode as another directory entry
 - can't span filesystems
 - can't create hard links to non-existent file
 - can't reference directories
 - do not occupy storage space (i.e. blocks)
 - `ln [option]... target link_name`



- `ls -i`

Filesystem symbolic links

- a file that references another file via path and name
 - can reference directories
 - can span filesystems
 - can reference non-existent files
 - `ln -s target link_name`
 - occupy space



- `symlink` / soft link

df, du, stat

- df Report disk space usage per filesystem
 - -h human readable output
 - -i list inode information instead of block usage
 - -T include filesystem type
 - --si use powers of 1000 instead of 1024
 - -P use the POSIX output format
- du Report disk usage per file and directory
 - -h human readable sizes
 - -s summarize, only display total for each argument
 - -x do not include files on a different filesystem
 - --si use powers of 1000 instead of 1024
- stat display file or file system status
 - -L follow links
 - -c --format

File extensions and content

- file extensions are just part of the file name
- some applications may care about extensions
- `file` - reports the type of file by examining the file contents
- `/usr/share/file/magic.mgc`

Displaying text files

- `cat` - concatenate files and print on the standard output
- `more`
- `less`
- `head`
- `tail`
 - `tail -f`
- `-n`

Displaying binary files

- displaying raw binary data may corrupt the display terminal
- `strings` - displays ASCII text inside binary files
- `xxd` - displays HEX and ASCII dump of file
- `clear`

xargs

- build and execute command lines from standard input
- `xargs [options] [command [initial-arguments]]`
- reads items from the standard input
 - delimited by blanks or newlines
- executes the command (`/bin/echo`)
 - one or more times
 - with any initial-arguments
 - followed by items read from standard input
- `-0, --null`
- `-I`
- `-n`

Searching the filesystem

- *machine-readable*
- `find [options] [starting-point] [expression]`
 - global/positional options
 - tests
 - operators
 - `-o, -a` (default)
 - actions
 - `-print` vs `-print0` vs `-printf`
 - `-ls`
 - `-exec`
- `find /foo -name bar -print`

Archiving & compressing

- archiving
 - tar
 - cpio
- compressing
 - compress
 - gzip
 - bzip2
 - lzma
 - xz

Archives with tar

- tar
 - manipulates .tar files (*tarballs*)
 - used for backup and transfer of files
 - creates, extracts or lists the contents of tarballs
 - c, x, t, f, v
 - traditional vs. UNIX-style vs. GNU-style usage
 - `tar cvf foo.tar ./foo/*`
 - GNU tar supports built-in compression methods
 - `-a, --auto-compress`
 - `-J, --xz`
- .tar (*tarball*)
 - records file and directory structure
 - includes metadata about the file: date, timestamps, ownership, permissions, etc.

XZ Utils

- `xz`
- `.xz`
- `unxz` / `xzcat` / `xz -d`
- compression format of choice